

Correlated Exploration for Cooperative MARL

via a Gaussian Copula

Project 9 — Reinforcement Learning

The problem: coordinated exploration

In cooperative MARL the team is often rewarded **only** when agents act *together* (e.g. both reach a goal in the same step).

- Standard ϵ -greedy explores each agent **independently**
- Probability all n agents explore the matching joint action $\propto \epsilon^n$
- $\rightarrow$ decays **exponentially** in the number of agents

Hypothesis: coupling agents' exploration — so similar agents tend to act alike — should accelerate learning on coordination-bottlenecked tasks.

Method: Gaussian copula

Keep the ε -greedy template, but **correlate** the random decisions:

1. Build a correlation matrix R (PSD, unit diagonal)
2. Cholesky $R = LL^\top$
3. Sample $z \sim \mathcal{N}(0, I)$, set $y = Lz \sim \mathcal{N}(0, R)$
4. $u_i = \Phi(y_i) \rightarrow$ uniform marginals, **correlated** across agents

Agent i explores iff $u_i < \varepsilon$. Independent ε -greedy is exactly the special case $R = I$.

Where does the correlation come from?

R = **cosine similarity** of per-agent feature vectors → a Gram matrix, automatically PSD with unit diagonal (+ small jitter).

Three similarity sources compared:

- **obs** — the agent's observation
- **q_values** — its current Q-values
- **hidden** — its Q-network backbone activations

Role-agnostic: nothing tells the method which agents share a role — similar agents simply get correlated automatically.


Experimental setup

Environments

- **Hallway** (MAVEN) — 2 agents, separate corridors, reward only if both reach 0
- **Level-Based Foraging** — 2 agents must load food *together* (sparse)
- **SMAX** `2s3z` (StarCraft-like, JAX) — 5 agents, 2 stalkers + 3 zealots


Algorithms: VDN (sum) and QMIX (monotonic mixer), parameter-shared Q-network, DQN-style training, **5 seeds** each.

Result 1 — Hallway: faster early learning




Correlated- **obs** early success (ep 50–100): **+64%** (VDN), **+94%** (QMIX) over independent. All converge eventually → the benefit is a **speed-up**.

Result 2 — LBF: the mechanism works even where learning fails




Catastrophic forgetting collapses all methods, **but** correlated reaches a **2.6–2.9× higher peak** — it finds the coordinated successes independent exploration misses.

Result 3 — SMAX 2s3z: scales to GPU



All methods reach ~**100% win rate**. Correlated reaches 50% win ~**11% (VDN)** / ~**15% (QMIX) faster** than independent. (*Getting a from-scratch JAX learner to converge required Double DQN + soft targets + state LayerNorm.*)

Result 4 — what the correlation matrix reveals



obs keeps role structure (stalker pair 0.80 vs. 0.25 to zealots) — the observation always encodes unit type. **q_values** / **hidden** homogenise (~0.95+) as the policy converges → why **obs** is the reliable source.

Conclusions

- Correlated ε -greedy via a Gaussian copula is a **simple, drop-in** change to exploration
- **Accelerates** learning where coordination is the bottleneck (Hallway +64–94%, SMAX 11–15% faster)
- **Improves exploration quality** even where value learning fails (LBF 2.6–2.9×)
- **obs** is the most reliable similarity source
- The correlation matrix **recovers role structure** automatically

Thank you — questions?